

Deep Generative Models: Diffusion Models

Fall Semester 2025

René Vidal

Director of the Center for Innovation in Data Engineering and Science (IDEAS)
Rachleff University Professor, University of Pennsylvania
Amazon Scholar & Chief Scientist at NORCE



Outline


- Markov Hierarchical Variational Auto Encoders (MHVAEs)
 - Autoregressive Encoder and Autoregressive Decoder of an MHVAE
 - Derivation of the ELBO of an MHVAE
- **Diffusion Models: MHVAEs with a Linear Gaussian Autoregressive Latent Space**
 - Forward Diffusion Process
 - Reverse Diffusion Process
 - ELBO for Diffusion Models as a particular case of the ELBO for MHVAEs
- Implementation Details: UNet architecture, Training and Sampling Strategies
- Application of Diffusion Models
 - Text-Conditioned Diffusion Model: Stable Diffusion
 - Multimodal Control for Consistent Synthesis: ControlNet
 - Image Editing: DDIM, P2P

Diffusion Model as MHVAEs with Gaussian Latents

- A Diffusion Model is an MHVAE where the latent variables $\mathbf{x}_{1:T}$ have the same dimension as the data x_0 , and the encoder $q_\phi(x_{1:T} \mid x_0) = \prod_{t=1}^T q_\phi(x_t \mid x_{t-1})$ is not learned, but it is pre-specified as a linear Gaussian model

$$q_\phi(x_t \mid x_{t-1}) = \mathcal{N}(x_t; \sqrt{\alpha_t}x_{t-1}, (1 - \alpha_t)I)$$
$$x_t = \sqrt{\alpha_t}x_{t-1} + \sqrt{1 - \alpha_t} \epsilon_t, \quad \epsilon_t \sim \mathcal{N}(\epsilon_t; 0, I)$$

- The parameter α_t is chosen such that $x_T \sim \mathcal{N}(x_T; 0, I)$ is a standard Gaussian



The Forward Process of Diffusion Model

- Consider the formulation of a single noising step:

$$x_t = \sqrt{\alpha_t} x_{t-1} + \sqrt{1 - \alpha_t} \epsilon_t, \quad \epsilon_t \sim \mathcal{N}(\epsilon_t; 0, I).$$

- Note $x_1 \mid x_0$ and $x_t \mid x_{t-1}$ are Gaussian, hence $x_t \mid x_0 \sim \mathcal{N}(x_t; \mu_q(x_0), \Sigma_q(x_0))$.
- We can compute $\mu_q(x_0) = \mathbb{E}[x_t \mid x_0]$, recursively as follows:

$$\begin{aligned}\mathbb{E}[x_t \mid x_0] &= \mathbb{E}\left[\sqrt{\alpha_t} x_{t-1} + \sqrt{1 - \alpha_t} \epsilon_t \mid x_0\right] \\ &= \sqrt{\alpha_t} \mathbb{E}[x_{t-1} \mid x_0] + \sqrt{1 - \alpha_t} \mathbb{E}[\epsilon_t] \quad \text{0} \\ &= \sqrt{\alpha_t} \mathbb{E}[x_{t-1} \mid x_0] \\ &= \sqrt{\alpha_t} \sqrt{\alpha_{t-1}} \mathbb{E}[x_{t-2} \mid x_0] \\ &= \sqrt{\alpha_t} \sqrt{\alpha_{t-1}} \cdots \sqrt{\alpha_1} x_0 \\ &= \sqrt{\alpha_t} x_0\end{aligned}$$

$$\bar{\alpha}_t = \prod_{i=1}^t \alpha_i$$

The Forward Process of Diffusion Model

- We can compute $\Sigma_q(x_0) = \text{Var}[x_t | x_0]$, recursively as follows:

$$\begin{aligned}\text{Var}(x_t | x_0) &= \text{Var}(\sqrt{\alpha_t}x_{t-1} + \sqrt{1 - \alpha_t}\epsilon_t | x_0) \\ &= \alpha_t \text{Var}(x_{t-1} | x_0) + (1 - \alpha_t) \text{Var}(\epsilon_t) \\ &= \alpha_t \text{Var}(x_{t-1} | x_0) + (1 - \alpha_t) I\end{aligned}$$

- That is:

$$\begin{aligned}\text{Var}(x_t | x_0) &= \alpha_t [\alpha_{t-1} \text{Var}(x_{t-2} | x_0) + (1 - \alpha_{t-1}) I] + (1 - \alpha_t) I \\ &= \alpha_t \alpha_{t-1} \text{Var}(x_{t-2} | x_0) + (1 - \alpha_t \alpha_{t-1}) I \\ &= \dots \\ &= \alpha_t \alpha_{t-1} \dots \alpha_1 \text{Var}(x_0 | x_0) + \left(1 - \prod_{i=1}^t \alpha_i\right) I \\ &= \left(1 - \prod_{i=1}^t \alpha_i\right) I = (1 - \bar{\alpha}_t) I\end{aligned}$$

$$\boxed{\bar{\alpha}_t = \prod_{i=1}^t \alpha_i}$$

The Forward Process of Diffusion Model

- We have shown that x_t is a linear Gaussian transformation of x_0 with scheduled randomness (controlled by $\bar{\alpha}_t$) drawn from a standard normal distribution, i.e.,

$$x_t \mid x_0 \sim \mathcal{N}\left(x_t; \sqrt{\bar{\alpha}_t}x_0, (1 - \bar{\alpha}_t)I\right)$$

- Therefore, given x_0 , we can sample x_t directly without having to generate all x_t 's:

$$x_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t} \bar{\epsilon}_t, \quad \bar{\epsilon}_t \sim \mathcal{N}(\bar{\epsilon}_t; 0, I)$$

- Moreover, we can also generate x_0 from x_t as

$$x_0 = (x_t - \sqrt{1 - \bar{\alpha}_t} \bar{\epsilon}_t) / \sqrt{\bar{\alpha}_t}, \quad \bar{\epsilon}_t \sim \mathcal{N}(\bar{\epsilon}_t; 0, I)$$

- This suggests we can reverse the noising process. However, exact reversal requires knowing the exact $\bar{\epsilon}_t$. The reverse diffusion process is designed to predict the noise $\bar{\epsilon}_t$ that needs to be added to x_t to generate x_0 .

The Reverse Diffusion Process

- We have designed a forward diffusion process $q_\phi(x_t | x_{t-1})$ that
 - At each step adds Gaussian noise to the input until it becomes pure noise
 - Allows us to sample $x_t | x_0$ without having to compute x_t recursively
 - Allows us to sample $x_0 | x_t$ without having to compute x_t recursively
- We now need to design a reverse diffusion process $p_\theta(x_{t-1} | x_t)$ that makes the calculation of the ELBO easy. We do this by
 - Understanding the structure of $q_\phi(x_t | x_{t-1})$
 - Making $p_\theta(x_{t-1} | x_t)$ match that structure
- Recall the ELBO is given by:

$$= \underbrace{\mathbb{E}_{q_\phi(x_1|x_0)}[\log p_\theta(x_0 | x_1)]}_{\text{reconstruction term}} - \underbrace{D_{\text{KL}}(q_\phi(x_T | x_0) || p_\theta(x_T))}_{\text{prior matching term}} - \sum_{t=2}^T \underbrace{\mathbb{E}_{q_\phi(x_t|x_0)}[D_{\text{KL}}(q_\phi(x_{t-1} | x_t, x_0) || p_\theta(x_{t-1} | x_t))]}_{\text{score matching term}}$$

ELBO for Diffusion Model: Score Matching Term

- To compute the third term, we need

$$q(x_{t-1} | x_t, x_0) = \frac{q(x_t | x_{t-1}, x_0) q(x_{t-1} | x_0)}{q(x_t | x_0)} = \frac{\mathcal{N}(x_t; \sqrt{\alpha_t} x_{t-1}, (1 - \alpha_t)I) \mathcal{N}(x_{t-1}; \sqrt{\alpha_{t-1}} x_0, (1 - \alpha_{t-1})I)}{\mathcal{N}(x_t; \sqrt{\alpha_t} x_0, (1 - \alpha_t)I)}$$

$$\propto \exp\left(-\frac{\|x_t - \sqrt{\alpha_t} x_{t-1}\|^2}{2(1 - \alpha_t)}\right) \cdot \exp\left(-\frac{\|x_{t-1} - \sqrt{\alpha_{t-1}} x_0\|^2}{2(1 - \alpha_{t-1})}\right) \propto \exp\left(-\frac{\alpha_t \|x_{t-1} - \frac{x_t}{\sqrt{\alpha_t}}\|^2}{2(1 - \alpha_t)}\right) \cdot \exp\left(-\frac{\|x_{t-1} - \sqrt{\alpha_{t-1}} x_0\|^2}{2(1 - \alpha_{t-1})}\right)$$

- Applying the product rule, we get $q(x_{t-1} | x_t, x_0) = \mathcal{N}(x_{t-1}; \mu_q(x_0), \Sigma_q)$, where

$$\Sigma_q := \text{Cov}(x_{t-1} | x_t, x_0) = \left(\frac{\alpha_t}{1 - \alpha_t} I + \frac{1}{1 - \alpha_{t-1}} I \right)^{-1} = \frac{(1 - \alpha_t)(1 - \alpha_{t-1})}{1 - \alpha_t} I$$

$$\mu_q(x_t, x_0) := \mathbb{E}(x_{t-1} | x_t, x_0) = \Sigma_q \left(\frac{\alpha_t}{1 - \alpha_t} I \frac{x_t}{\sqrt{\alpha_t}} + \frac{1}{1 - \alpha_{t-1}} I \sqrt{\alpha_{t-1}} x_0 \right)$$

$$= \frac{(1 - \alpha_t)(1 - \alpha_{t-1})}{1 - \alpha_t} I \left(\frac{\sqrt{\alpha_t}}{1 - \alpha_t} x_t + \frac{\sqrt{\alpha_{t-1}}}{1 - \alpha_{t-1}} x_0 \right) = \frac{\sqrt{\alpha_t}(1 - \alpha_{t-1})x_t + \sqrt{\alpha_{t-1}}(1 - \alpha_t)x_0}{1 - \alpha_t}$$

$q(x_t x_{t-1}) = \mathcal{N}(x_t; \sqrt{\alpha_t} x_{t-1}, (1 - \alpha_t)I)$
$q(x_t x_0) = \mathcal{N}(x_t; \sqrt{\alpha_t} x_0, (1 - \alpha_t)I)$

$\mathcal{N}(x; \mu_1, \Sigma_1) \mathcal{N}(x; \mu_2, \Sigma_2) \propto \mathcal{N}(x; \bar{\mu}, \bar{\Sigma})$
$\bar{\mu} = \bar{\Sigma}(\Sigma_1^{-1}\mu_1 + \Sigma_2^{-1}\mu_2)$, $\bar{\Sigma} = (\Sigma_1^{-1} + \Sigma_2^{-1})^{-1}$

$\mu_1 = \frac{x_t}{\sqrt{\alpha_t}}$, $\Sigma_1 = \frac{1 - \alpha_t}{\alpha_t} I$,
$\mu_2 = \sqrt{\alpha_{t-1}} x_0$, $\Sigma_2 = (1 - \alpha_{t-1}) I$

ELBO for Diffusion Model: Decoder Matches Encoder

- Recall KL divergence for Gaussians

$$D_{\text{KL}} \left(\mathcal{N}(x; \mu_x, \Sigma_x) \parallel \mathcal{N}(y; \mu_y, \Sigma_y) \right) = \frac{1}{2} \left[\log \frac{|\Sigma_y|}{|\Sigma_x|} - d + \text{tr}(\Sigma_y^{-1} \Sigma_x) + (\mu_y - \mu_x)^T \Sigma_y^{-1} (\mu_y - \mu_x) \right]$$

~~$|\Sigma_y|$~~ ~~$\text{tr}(\Sigma_y^{-1} \Sigma_x)$~~

- Choosing mean of $p_\theta(x_{t-1} | x_t)$ to match form of mean of $q(x_{t-1} | x_t, x_0)$

$$\mu_q(x_t, x_0) = \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})x_t + \sqrt{\bar{\alpha}_{t-1}}(1 - \alpha_t)x_0}{1 - \bar{\alpha}_t} \Rightarrow \mu_\theta(x_t, t) = \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})x_t + \sqrt{\bar{\alpha}_{t-1}}(1 - \alpha_t)\widehat{x_\theta}(x_t, t)}{1 - \bar{\alpha}_t}$$

- Choosing variance of $p_\theta(x_{t-1} | x_t)$ to match exactly variance of $q(x_{t-1} | x_t, x_0)$

$$\Sigma_q = \frac{(1 - \alpha_t)(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} I \Rightarrow \Sigma_\theta = \frac{(1 - \alpha_t)(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} I$$

- The ELBO reduces to:

$$D_{\text{KL}}(q(x_{t-1} | x_t, x_0) \parallel p_\theta(x_{t-1} | x_t)) = D_{\text{KL}} \left(\mathcal{N}(x_{t-1}; \mu_q(x_t, x_0), \Sigma_q) \parallel \mathcal{N}(x_{t-1}; \mu_\theta(x_t, t), \Sigma_\theta) \right)$$

$$= \frac{1}{2\sigma_q^2(t)} [\|\mu_\theta - \mu_q\|_2^2] = \frac{1}{2\sigma_q^2(t)} \frac{\bar{\alpha}_{t-1}(1 - \alpha_t)^2}{(1 - \bar{\alpha}_t)^2} [\|\widehat{x_\theta}(x_t, t) - x_0\|_2^2]$$

$\sigma_q^2(t) = \frac{(1 - \alpha_t)(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t}$

Reparameterization as an Alternative Form for ELBO

- Plugging $x_0 = \frac{x_t - \sqrt{1-\bar{\alpha}_t} \hat{\epsilon}_\theta(x_t, t)}{\sqrt{\bar{\alpha}_t}}$ into the denoising transition mean $\mu_q(x_t, x_0)$, we obtain:

$$\begin{aligned}
 \mu_q(x_t, x_0) &= \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})x_t + \sqrt{\bar{\alpha}_{t-1}}(1 - \alpha_t)x_0}{1 - \bar{\alpha}_t} \\
 &= \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})x_t + \sqrt{\bar{\alpha}_{t-1}}(1 - \alpha_t)\frac{x_t - \sqrt{1 - \bar{\alpha}_t} \hat{\epsilon}_\theta(x_t, t)}{\sqrt{\bar{\alpha}_t}}}{1 - \bar{\alpha}_t} \\
 &= \frac{1 - \bar{\alpha}_t}{(1 - \bar{\alpha}_t)\sqrt{\alpha_t}} x_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}\sqrt{\alpha_t}} \hat{\epsilon}_\theta(x_t, t) \\
 &= \frac{1}{\sqrt{\alpha_t}} x_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}\sqrt{\alpha_t}} \hat{\epsilon}_\theta(x_t, t)
 \end{aligned}$$

- Choosing the mean and variance of p to match the mean and variance of q :

$$\mu_\theta(x_t, t) = \frac{1}{\sqrt{\alpha_t}} x_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}\sqrt{\alpha_t}} \hat{\epsilon}_\theta(x_t, t), \quad \Sigma_\theta(t) = \frac{(1 - \alpha_t)(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} I$$

The Reverse Diffusion Process for DDPM

- Putting it all together, the reverse diffusion process is given by:

$$p_{\theta}(x_{t-1} | x_t) \sim \mathcal{N}(x_{t-1}; \mu_{\theta}(x_t, t), \Sigma_{\theta}(t))$$

- In DDPM, $\Sigma_{\theta}(t)$ is not learned and is fixed to $\Sigma_{\theta}(t) = \beta_t I$ to match the noise added in the forward process.
- Therefore, we generate an image via the reverse diffusion process


$$x_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \hat{\epsilon}_{\theta}(x_t, t) \right) + \sqrt{\beta_t} \epsilon_t$$

where $x_T \sim \mathcal{N}(x_T; 0, I)$, $\epsilon_t \sim \mathcal{N}(\epsilon_t; 0, I)$.

Progressive Denoising or Direct Reconstruction?

- The model predicts the **noise** to be removed in each step by optimizing the **score matching term**. This reduces to minimizing the difference between the predicted noise and the ground-truth schedule noise:

$$\begin{aligned} & \underset{\theta}{\operatorname{argmin}} D_{\text{KL}}(q(x_{t-1} | x_t, x_0) \| p_{\theta}(x_{t-1} | x_t)) \\ &= \underset{\theta}{\operatorname{argmin}} D_{\text{KL}}\left(\mathcal{N}\left(x_{t-1}; \mu_q, \Sigma_q(t)\right) \| \mathcal{N}\left(x_{t-1}; \mu_{\theta}, \Sigma_q(t)\right)\right) \\ &= \underset{\theta}{\operatorname{argmin}} \frac{1}{2\sigma_q^2(t)} \left[\left\| \frac{1}{\sqrt{\alpha_t}} x_t - \frac{1-\alpha_t}{\sqrt{1-\bar{\alpha}_t} \sqrt{\alpha_t}} \hat{\epsilon}_{\theta}(x_t, t) - \frac{1}{\sqrt{\alpha_t}} x_t + \frac{1-\alpha_t}{\sqrt{1-\bar{\alpha}_t} \sqrt{\alpha_t}} \epsilon_0 \right\|_2^2 \right] \\ &= \underset{\theta}{\operatorname{argmin}} \frac{1}{2\sigma_q^2(t)} \frac{(1-\alpha_t)^2}{(1-\bar{\alpha}_t)\alpha_t} [\|\epsilon_0 - \hat{\epsilon}_{\theta}(x_t, t)\|_2^2] \end{aligned}$$



- Predicting x_0 from a highly noisy x_t in one step is complex, as the signal is dominated by significant noise for large t .
- Predicting the noise at each step and refining x_t towards x_0 makes learning more manageable (e.g., it converges faster or it requires a smaller network capacity).